

Deep Learning

Understanding Data

```
train_df = pd.read_csv('train.csv')
test_df = pd.read_csv('test.csv')
store_df = pd.read_csv('store.csv')
```

A	B	C	D	E	F	G	H	I	
Store	DayOfWeek	Date	Sales	Customers	Open	Promo	StateHoliday	SchoolHoliday	
1	5	7/31/2015	5263	555	1	1	0	1	
2	5	7/31/2015	6064	625	1	1	0	1	
3	5	7/31/2015	8314	821	1	1	0	1	
4	5	7/31/2015	13995	1498	1	1	0	1	

Store	StoreType	Assortment	Competitor	Competitor	Competitor	Promo2	Promo2Size	Promo2Size	PromoInterval
1	c	a	1270	9	2008	0			
2	a	a	570	11	2007	1	13	2010	Jan, Apr, Jul, Oct
3	a	a	14130	12	2006	1	14	2011	Jan, Apr, Jul, Oct
4	c	c	620	9	2009	0			

Pre-Processing

Removing the column that are not in the test

Data frame



```
#Dropping out Id coloumn to make it smoother
test_df.drop(['Id'],axis=1,inplace=True)
train_df.drop(["Customers"],axis=1,inplace=True)
```

Checking if there are any missing values
in Train.CSV

```
# Checking for missing values in the training dataset
missing_values = train_df.isnull().sum()
missing_percentage = (missing_values / len(train_df)) * 100
missing_data_info = pd.DataFrame({'Missing_Values': missing_values, 'Missing_Percentage': missing_percentage})
missing_data_info = missing_data_info.sort_values(by='Missing_Percentage', ascending=False)
print("Columns with the highest percentage of missing values:")
print(missing_data_info.head())
```

Columns with the highest percentage of missing values:

	Missing_Values	Missing_Percentage
Store	0	0.0
DayOfWeek	0	0.0
Date	0	0.0
Sales	0	0.0
Open	0	0.0

Pre-Processing

Checking if there are any missing values
in test.CSV

```
Columns with the highest percentage of missing values:
```

	Missing_Values	Missing_Percentage
Open	11	0.026772
Store	0	0.000000
DayOfWeek	0	0.000000
Date	0	0.000000
Promo	0	0.000000

```
test_df.fillna(0, inplace=True)
```

```
Columns with the highest percentage of missing values:
```

	Missing_Values	Missing_Percentage
Store	0	0.0
DayOfWeek	0	0.0
Date	0	0.0
Open	0	0.0
Promo	0	0.0

Pre-Processing

Merging The data with the store data provided

```
train_store = pd.merge(train_df, store_df, on='Store')  
test_store = pd.merge(test_df, store_df, on='Store')  
train_store.head()
```

Checking if there are any missing values
in Store.CSV

```
missing_values = train_store.isnull().sum()
```

Columns with the highest percentage of missing values:

	Missing_Values	Missing_Percentage
PromoInterval	508031	49.943620
Promo2SinceYear	508031	49.943620
Promo2SinceWeek	508031	49.943620
CompetitionOpenSinceYear	323348	31.787764
CompetitionOpenSinceMonth	323348	31.787764

Pre-Processing

Filling the Missing Values in the Store.CSV

After testing

- median
- mean
- zeros

```
train_store['CompetitionDistance'].fillna(train_store['CompetitionDistance'].median(), inplace=True)
test_store['CompetitionDistance'].fillna(test_store['CompetitionDistance'].median(), inplace=True)
```

```
train_store['PromoInterval'].fillna(0, inplace=True)
test_store['PromoInterval'].fillna(0, inplace=True)
```

```
train_store['CompetitionOpenSinceYear'].fillna(0, inplace=True)
train_store['CompetitionOpenSinceMonth'].fillna(0, inplace=True)
test_store['CompetitionOpenSinceYear'].fillna(0, inplace=True)
test_store['CompetitionOpenSinceMonth'].fillna(0, inplace=True)
```

```
train_store['Promo2SinceYear'].fillna(0, inplace=True)
train_store['Promo2SinceWeek'].fillna(0, inplace=True)
test_store['Promo2SinceYear'].fillna(0, inplace=True)
test_store['Promo2SinceWeek'].fillna(0, inplace=True)
```

```
print("Closed stores with sales are ",((train_store["Open"]==0) & (train_store["Sales"]>0)).sum())
print("Open stores with zero sales are",((train_store["Open"]==1) & (train_store["Sales"]<=0)).sum())
print("negative sales are ",(train_store["Sales"]<0).sum())
```

```
Closed stores with sales are 0
Open stores with zero sales are 54
negative sales are 0
```

```
train_store = train_store[~((train_store["Open"] == 1) & (train_store["Sales"] <= 0))]
```

Removing illogical Cases:

Pre-Processing

Removing Outliers

```
z_scores = np.abs(zscore(numeric))
outliers = (z_scores > 3)
print(f"Number of rows before removing outliers: {train_store.shape[0]}")
train_store = train_store[~np.any(outliers, axis=1)]
print(f"Number of rows after removing outliers: {train_store.shape[0]}")
```

```
Number of rows before removing outliers: 1017155
Number of rows after removing outliers: 986781
```

Encoding

Encoding Categorical Data

```
train_store['Date'] = pd.to_datetime(train_df['Date'])
test_store['Date'] = pd.to_datetime(test_df['Date'])

train_store['Week'] = pd.to_datetime(train_store['Date']).dt.isocalendar().week
train_store['DayOfYear'] = pd.to_datetime(train_store['Date']).dt.dayofyear
train_store['DayOfMonth'] = pd.to_datetime(train_store['Date']).dt.day
train_store['Year'] = pd.DatetimeIndex(train_store['Date']).year
train_store['Quarter'] = pd.DatetimeIndex(train_store['Date']).quarter
train_store['Month'] = pd.DatetimeIndex(train_store['Date']).month
```

```
train_store.drop(['Date'],axis=1,inplace=True)
```

Encoding and splitting Date data

Encoding

Encoding Categorical Data

```
#Manually encoding the stateholiday column before using pre made encoders:  
train_store['StateHoliday'].replace("a", 1, inplace=True)  
train_store['StateHoliday'].replace("b", 2, inplace=True)  
train_store['StateHoliday'].replace("c", 3, inplace=True)  
train_store['StateHoliday'].replace('0', 0, inplace=True)  
train_store['StateHoliday'].unique()  
test_store['StateHoliday'].replace("a", 1, inplace=True)  
test_store['StateHoliday'].replace("b", 2, inplace=True)  
test_store['StateHoliday'].replace("c", 3, inplace=True)  
test_store['StateHoliday'].replace('0', 0, inplace=True)  
test_store['StateHoliday'].unique()
```

Manually Encoding

Encoding

Encoding Categorical Data

```
from sklearn.preprocessing import LabelEncoder
#try one hot encoder maybe better
category_train_store = train_store.select_dtypes(exclude=[np.number]).columns.tolist()
print("Categorical features in train dataset:", category_train_store)

Categorical_test_store = test_store.select_dtypes(exclude=[np.number]).columns.tolist()
print("Categorical features in test dataset:", Categorical_test_store)

le = LabelEncoder()
for column in category_train_store:
    train_store[column] = le.fit_transform(train_store[column].astype(str))

for column in Categorical_test_store:
    test_store[column] = le.fit_transform(test_store[column].astype(str))

Categorical features in train dataset: ['StoreType', 'Assortment', 'PromoInterval']
Categorical features in test dataset: ['StoreType', 'Assortment', 'PromoInterval']
```

Using Label
Encoder

Tried Onehot encoder as well and get
dummies

Scaling

Normalizing Data Using MinMax

```
scaler = MinMaxScaler()
train_scaling=train_store.drop(["Sales"],axis=1,inplace=False)
scaled_x_train_store = pd.DataFrame(scaler.fit_transform(train_scaling),columns=train_scaling.columns)
scaled_test_store = scaler.fit_transform(test_store)
scaled_y_train_store = pd.DataFrame(scaler.fit_transform(pd.DataFrame(train_store["Sales"])),columns=["Sales"])
scaled_train_store=pd.concat([scaled_x_train_store, scaled_y_train_store], axis=1)
scaled_test_store = pd.DataFrame(scaled_test_store, columns=test_store.columns)
```

```
print(scaled_train_store.head())
```

	Store	DayOfWeek	Open	Promo	StateHoliday	SchoolHoliday	StoreType	\
0	0.000000	0.666667	1.0	1.0	0.0	1.0	0.666667	
1	0.000898	0.666667	1.0	1.0	0.0	1.0	0.000000	
2	0.001795	0.666667	1.0	1.0	0.0	1.0	0.000000	
3	0.002693	0.666667	1.0	1.0	0.0	1.0	0.666667	
4	0.004488	0.666667	1.0	1.0	0.0	1.0	0.000000	

	Assortment	CompetitionDistance	CompetitionOpenSinceMonth	...	\
0	0.0	0.045241	0.750000	...	
1	0.0	0.019906	0.916667	...	
2	0.0	0.510677	1.000000	...	
3	1.0	0.021716	0.750000	...	
4	0.0	0.010496	1.000000	...	


Splitting

Normalizing Data Using MinMax


```
x=scaled_train_store.drop(['Sales'],axis=1)
y=scaled_train_store['Sales']
print(y)
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=7)
```

Using 3 Different models


FCNN



ML algorithms




LSTM




FCNN

```
model = models.Sequential([
    layers.Input(shape=(x_train.shape[1],)),
    layers.Dense(256, activation='relu'),
    layers.BatchNormalization(),
    layers.Dropout(0.4),
    layers.Dense(128, activation='relu'),
    layers.BatchNormalization(),
    layers.Dropout(0.4), #s
    layers.Dense(64, activation='relu'),
    layers.BatchNormalization(),
    layers.Dropout(0.4),
    layers.Dense(32, activation='relu'),
    layers.Dense(1)
])
callbacks = [
    EarlyStopping(monitor='val_loss', patience=10, restore_best_weights=True),
    ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=1)
]
# Compile model
model.compile(optimizer='adam', loss='mean_squared_error', metrics=['mae'])
model.summary()

# Train the model
batch_size = 32
history = model.fit(x_train, y_train, validation_data=(x_test, y_test),
                    epochs=70, callbacks=callbacks, batch_size=batch_size)
```



FCNN



ML Algorithms

```
models = [
    ('Random Forest', RandomForestRegressor()),
    ('Ridge Regression', Ridge()),
    ('Support Vector Machine', SVR()),
    ('Neural Network', MLPRegressor())
]
# Hyperparameters for grid search (you can modify these for each model)
param_grid = {
    'Random Forest': {'n_estimators': [100, 300, 500]},
    'Ridge Regression': {'alpha': [0.1, 1.0, 10.0]},
    'Support Vector Machine': {'kernel': ['linear', 'rbf'], 'C': [1, 10]},
    'Neural Network': {'hidden_layer_sizes': [(50,), (100,)], 'alpha': [0.0001, 0.001]}
}

for model_name, model in models:
    grid_search = GridSearchCV(model, param_grid[model_name], scoring='neg_mean_squared_error', cv=5)
    grid_search.fit(X_train, y_train)
    y_pred = grid_search.predict(X_test)
    mse = mean_squared_error(y_test, y_pred)

    if mse < best_score:
        best_score = mse
        best_model = grid_search.best_estimator_
        best_model_name = model_name

    print(f"{model_name} - Mean Squared Error: {mse}")
print(f"Best Model: {best_model_name} - Best Mean Squared Error: {best_score}")
```


ML Algorithms

```
Random Forest - Mean Squared Error: 0.0005290457524361835
```

```
Ridge Regression - Mean Squared Error: 0.004696394952530777
```

LSTM


```
model = models.Sequential([
    layers.LSTM(64, return_sequences=True, input_shape=(x_train.shape[1], x_train.shape[2])),
    layers.Dropout(0.2),
    layers.LSTM(32),
    layers.Dropout(0.2),
    layers.Dense(1)
])
```

```
callbacks = [
    EarlyStopping(monitor='val_loss', patience=10, restore_best_weights=True),
    ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=1)
]
```

```
# Compile model
model.compile(optimizer='adam', loss='mean_squared_error', metrics=['mae'])
model.summary()
```

```
# Train the model
batch_size = 32
history = model.fit(x_train, y_train, validation_data=(x_test, y_test),
                    epochs=20, callbacks=callbacks, batch_size=batch_size)
```

LSTM



Testing the Model

6358/6358 ————— **19s** 3ms/step - loss: 6.9387e-04 - mae: 0.0168

Test Mean Absolute Error: 0.01686321571469307

1284/1284 ————— **3s** 2ms/step

Predicted Sales: [[6260.2036]

[8129.6177]

[10756.432]

...

[5460.703]

[21708.084]

[5102.3276]]

	ID	Sales
0	1	6260.203613
1	2	8129.617676
2	3	10756.431641
3	4	7594.935547
4	5	7585.748535
...
41083	41084	3141.491455
41084	41085	8225.785156
41085	41086	5460.703125
41086	41087	21708.083984
41087	41088	5102.327637

[41088 rows x 2 columns]

Thank You!